

drivy



How we use and operate Apache Airflow at Drivy

Eloïse GOMEZ & Faouz EL FASSI



drivy



Find a Drivy nearby

50,000 cars from locals
No subscription needed



Open it with the app

24/7 on-demand access
Prices from £29/day



Enjoy the drive!


Fully insured round trips
100 mi/day included



Data @ Drivy

- A 3 nodes Redshift WH (Snowflake soon)
- 6TB of data (including tracking)
- 2M users
- 8 employees (6 analysts + 2 engs.)

We are hiring in the data team.
Join us!



Summary

- How do we manage Airflow at Drivy?
- How do we use it?



How do we manage airflow at
Drivy?


Airflow Recap

“Airflow is a platform to programmatically author, schedule and monitor data pipelines.”


Airflow Components

- **Metadata Database:** stores information regarding tasks state.
- **Scheduler:** decides which tasks need to be executed.
- **Web server:** accepts HTTP requests and allows user to interact with Airflow.
- **Worker:** executes tasks.
- **A queuing service:** hold information about next task to execute.

Airflow Architecture



Local Executor



Celery Executor

Single-Node vs Multi-Nodes Architecture

Single-Node Architecture


- Easy to set up
- **Single point of failure**
- **Hard to scale**

Multi-Node Architecture

- Higher availability
- Dedicated workers for specific tasks
- Scaling horizontally

Airflow 6 months ago

- Single-node architecture
- 1 ec2 instance manually set up
- Crontable deployment tool

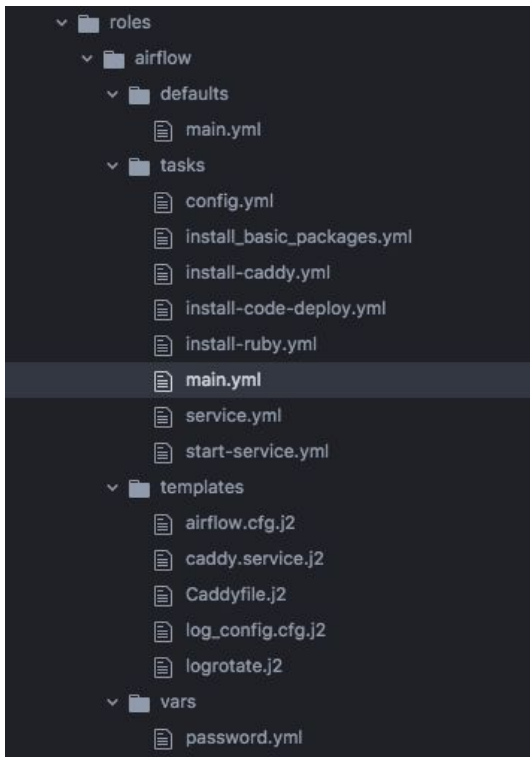


Airflow Architecture now

- Multi-node architecture
- 4 ec2 instances installed through ansible
- Code deploy as a deployment tool



Airflow Installation



```
- name: Install Airflow
  import_tasks: install_basic_packages.yml

- name: Set up caddy
  import_tasks: install-caddy.yml

- name: Install ruby on workers
  import_tasks: install-ruby.yml
  when: inventory_hostname in groups['airflow-workers']

- name: Configure Airflow
  import_tasks: config.yml

- name: Setting up Systemd
  import_tasks: service.yml

- name: Start services
  import_tasks: start-service.yml

- name: Deploy code-deploy agents
  import_tasks: install-code-deploy.yml
```

Airflow Installation

```
- name: Download airflow zip
  get_url:
    url: "https://codeload.github.com/apache/incubator-airflow/zip/{{ airflow_version }}"
    dest: /tmp/airflow.zip

- name : Create Airflow tmp directory
  file:
    path: /tmp/airflow
    state: directory
    mode: 0755

- name: Extract airflow.zip into /tmp/airflow
  unarchive:
    remote_src: yes
    src: /tmp/airflow.zip
    dest: /tmp/airflow
```

```
[airflow:children]
airflow-master
airflow-workers
```

```
[airflow-master]
```

```
master ansible_host=##.###.###.### ansible_port=22 ansible_user=ec2-user ansible_ssh_private_key_file=~/.secret_file
```


```
[airflow-workers]
```

```
worker-1 ansible_host=##.###.###.### ansible_port=22 ansible_user=ec2-user ansible_ssh_private_key_file=~/.secret_file
worker-2 ansible_host=##.###.###.### ansible_port=22 ansible_user=ec2-user ansible_ssh_private_key_file=~/.secret_file
worker-3 ansible_host=##.###.###.### ansible_port=22 ansible_user=ec2-user ansible_ssh_private_key_file=~/.secret_file
```

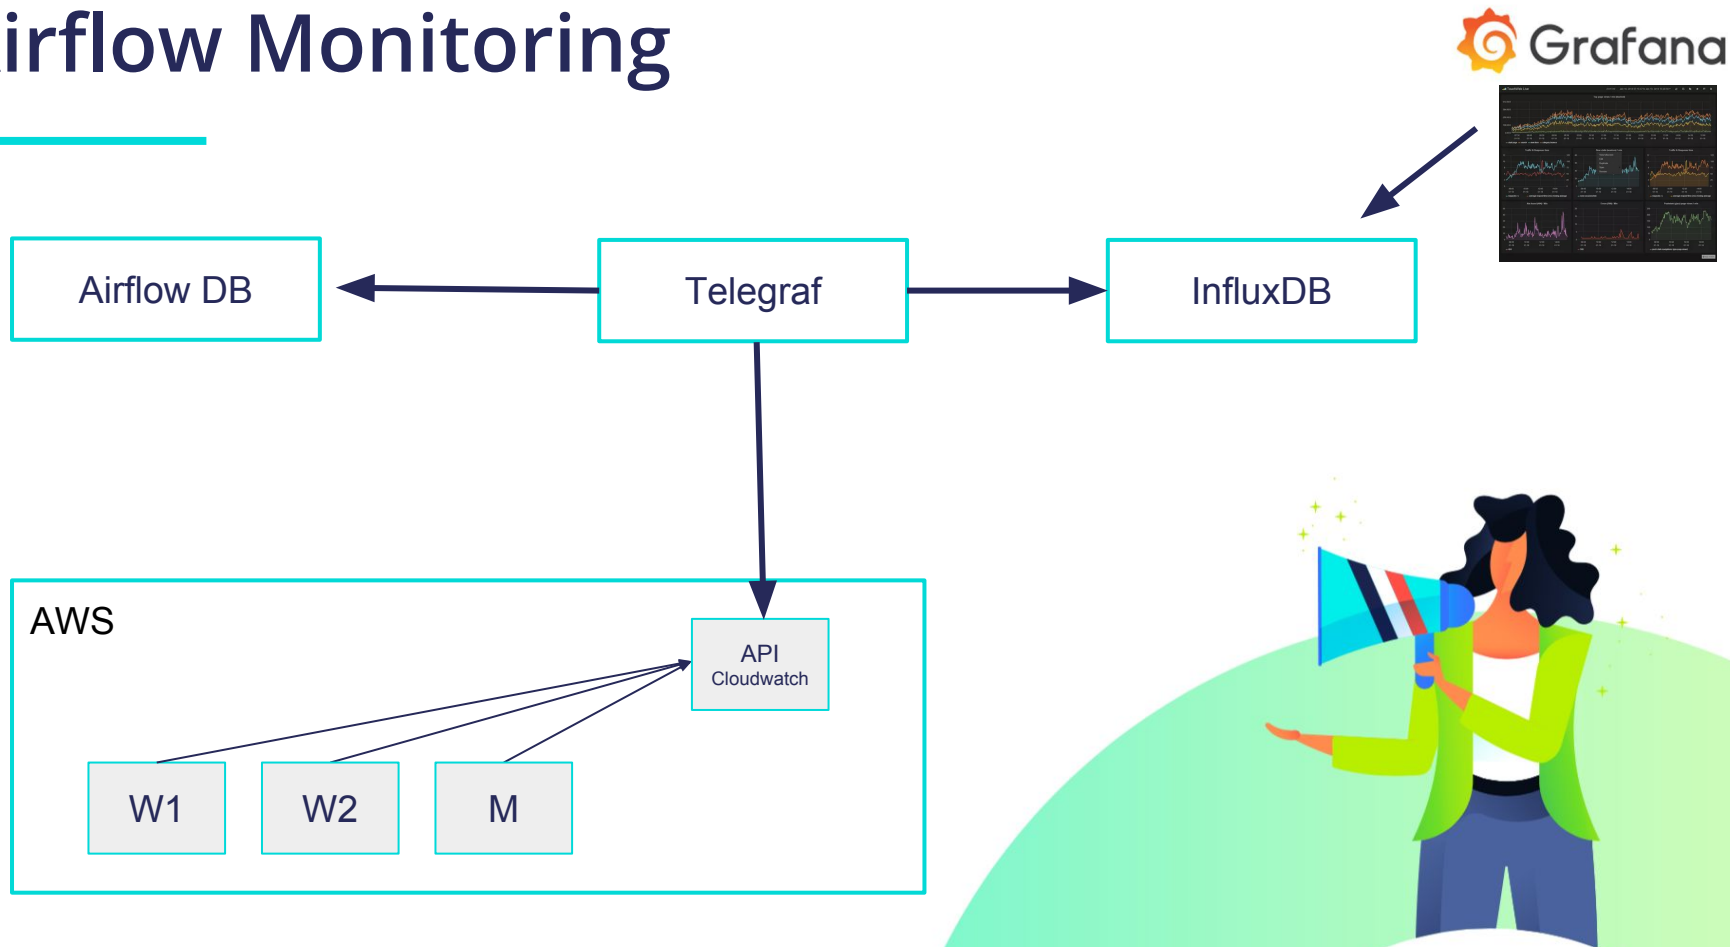


Airflow Migration pain points

- Migrate dag by dag to make sure they are celery executor compliant
- Airflow logs on S3
- Two dag folders
- Need to create symlinks for shared files (can't specify several dag folders)
- Maintain two Airflows



Airflow Monitoring



Airflow Monitoring

```
1 #!/usr/bin/env ruby
2
3 require_relative '../rb-helpers/mysql'
4 require_relative '../rb-helpers/exec'
5
6 dbs = ['airflow', 'airflow2']
7
8 dbs.each do |db|
9   mysql_client(db.to_sym).query(%Q{
10     SELECT state, count(1) count
11     FROM task_instance
12     WHERE end_date BETWEEN (now() - INTERVAL 1 MINUTE) AND now()
13     GROUP BY state
14   }, symbolize_keys: true).each do |row|
15     output(
16       prefix: db,
17       tags: { task_state: row[:state] },
18       values: {
19         count: row[:count]
20       }
21     )
22   end
23 end
```

```
[[inputs.exec]]
  interval = "1m"
  commands = [
    "bundle exec exec/mysql/airflow_tasks.rb"
  ]
  timeout = "10s"
  data_format = "influx"
```

Airflow Alerting



7:18 PM data-bot APP

[\[Alerting\] Not enough tasks processed by Airflow over the last 10 minutes](#)

airflow.sum

2

Grafana v4.6.2 | Jan 13th, 2018

7:26 PM data-bot APP

[\[OK\] Not enough tasks processed by Airflow over the last 10 minutes](#)

Grafana v4.6.2 | Jan 13th, 2018

Alert Config

Name: Not enough tasks processed by Airflow... Evaluate every: 600s

Conditions


WHEN sum () OF query (B, 10m, now) IS BELOW 10

+

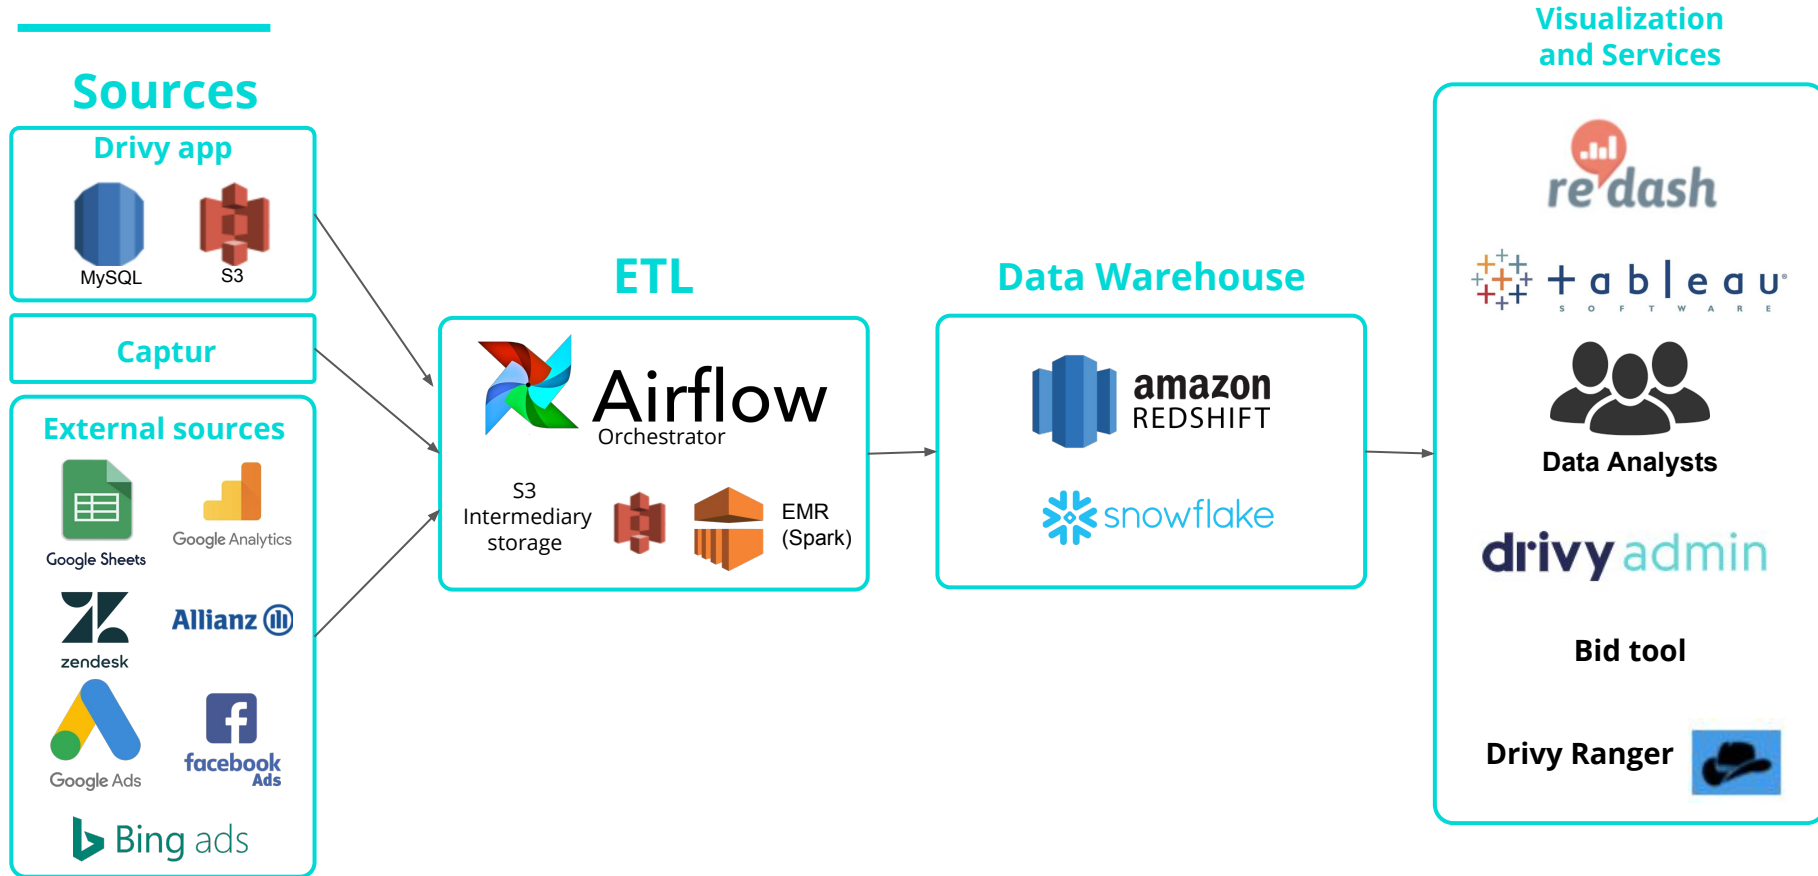
If no data or all values are null	SET STATE TO	No Data
If execution error or timeout	SET STATE TO	Alerting

Airflow tomorrow

- Improve Airflow deployment
- Test managed services
(Cloud Composer, Astronomer)
- Create a test environment



Infrastructure



How do we use Airflow at Drivy?


MySQL dump and enrichment

Three kinds of ingestions:

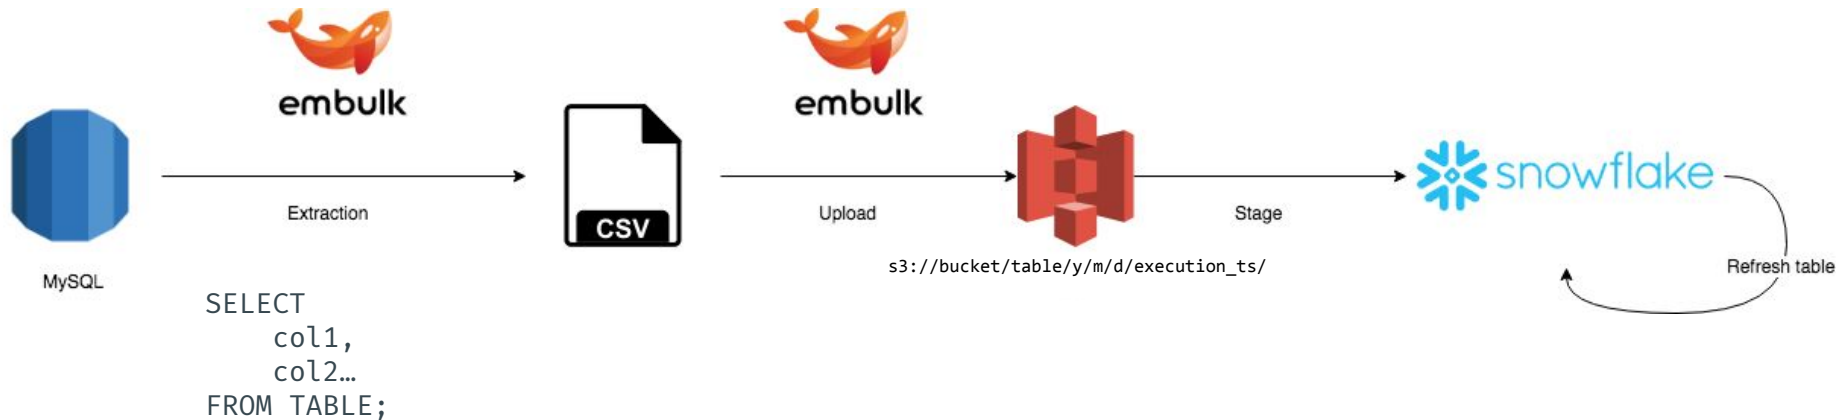
- Full dump: tables where records can be updated in the past without being timestamped, or having low volumes.
- Incremental append-only: tables where only new records are added.
- Incremental UPSERT: tables where new records are added and updated ones are timestamped.

Each of those can be **URGENT** (every 2h30 to 6h) or **NON URGENT** (every 6 to 12h)

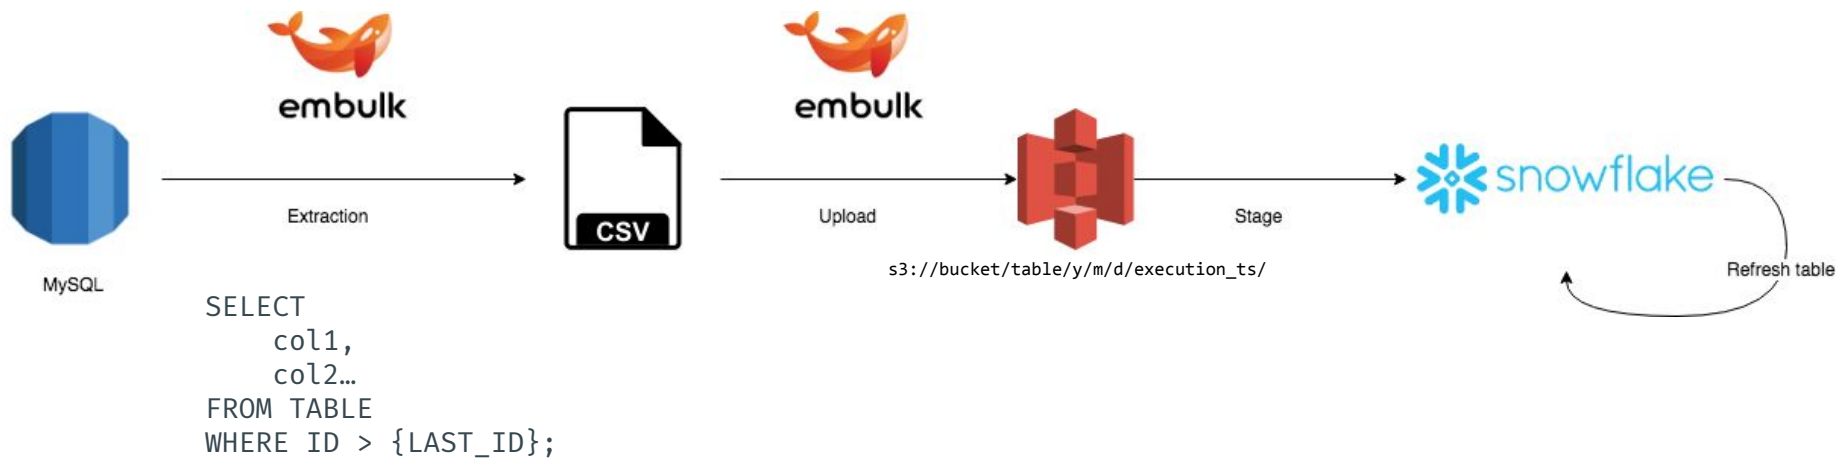
Full dump imports, the old way



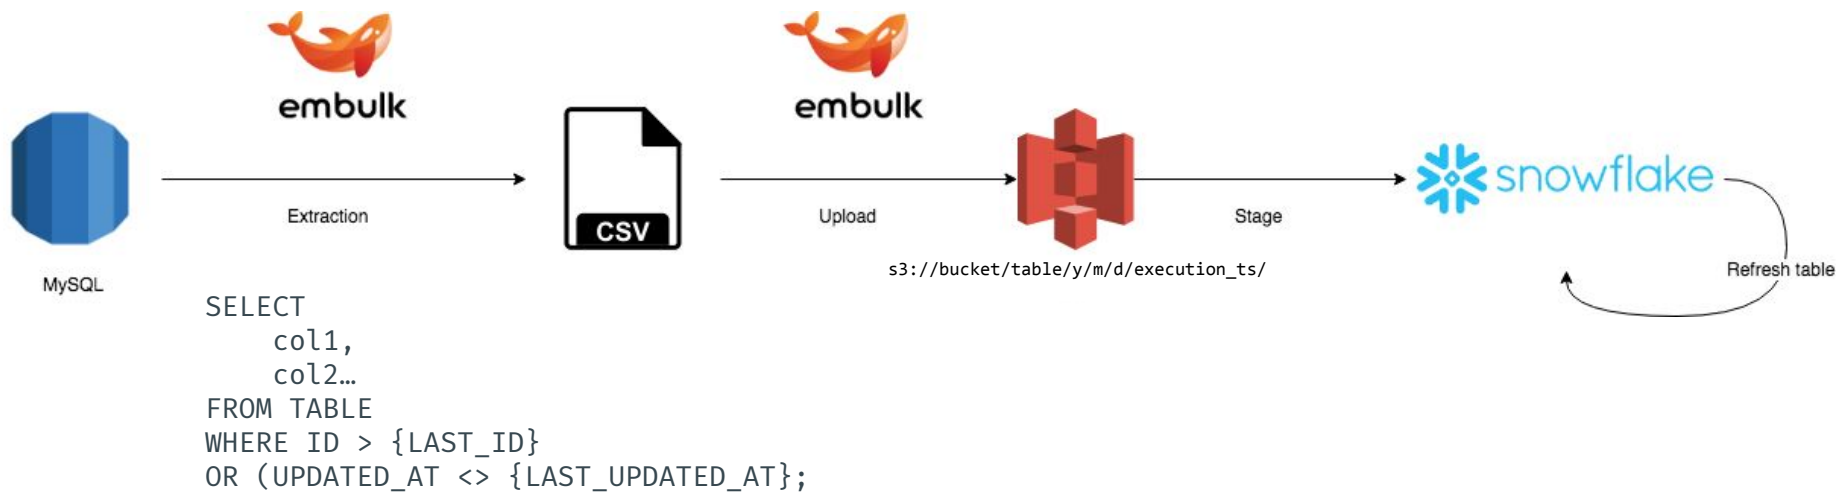
Full dump imports, the new way



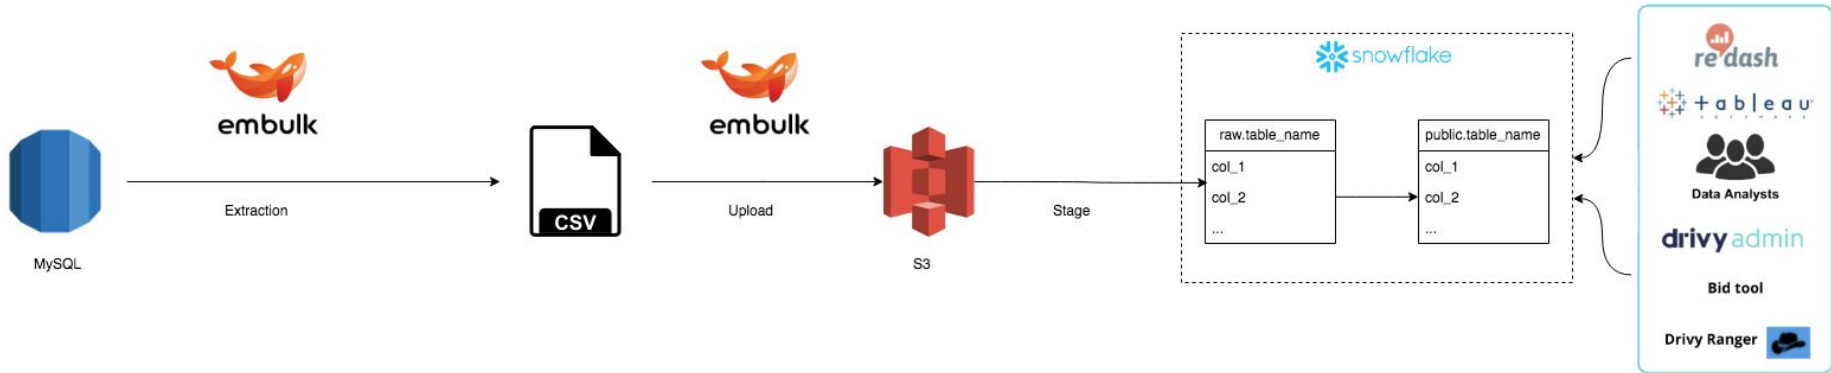
Incremental append



Incremental upsert



Raw / Transformed isolation



Captur

“Captur is our homemade cross-device tracking framework that mimics Segment architecture”

Captur - Usage

The goal of our backend engineers was to make sending events simple, for every developer.

```
tracking.event("cookie_banner_viewed")
```

← javascript

```
async_track(:event,  
  name: 'order_blocked',  
  order_id: @order.id,  
  reason: 'risky_picks',  
  client: :mobile  
)
```

← ruby

```
analyticsSendView(Tracking.Agreement.Mobile.Checkin.confirmation, objectId: rentalId)
```

← swift

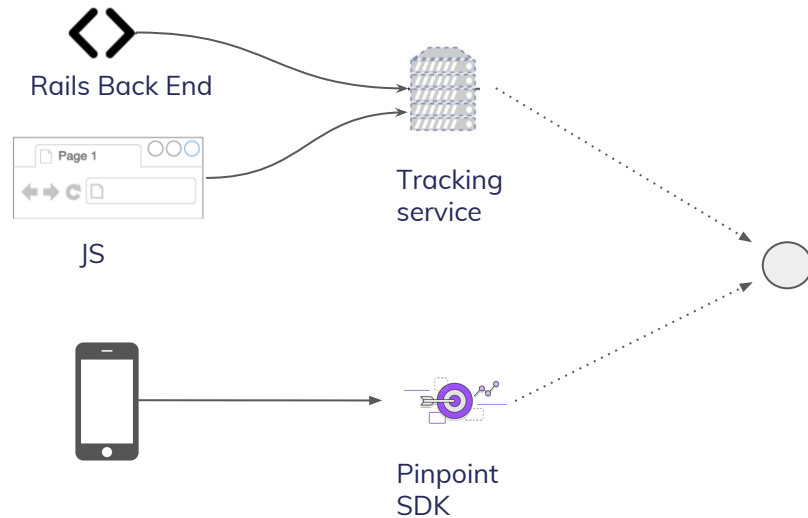
```
AnalyticsUtils.sendViewWithValues(AnalyticsEvents.DEEPLINK_FACEBOOK, values);
```

← java

Captur - Data collection

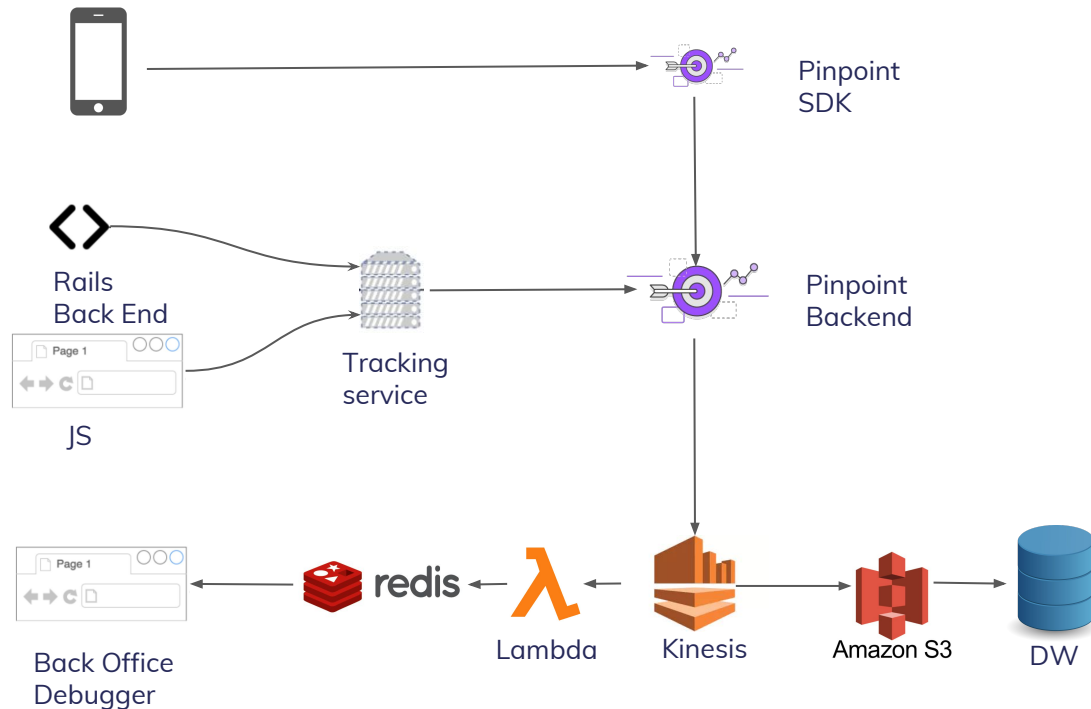
```
{
  "type": "event",
  "attributes": {
    "name": "car-preview",
    "car_id": "3082",
    "source": "car_card",
    "context": "instant_bookable",
    "title": "Renault Clio",
    "url": "https://staging.drivy.com/location-voiture/paris/r",
    "path": "/location-voiture/paris/renault-clio-3082",
    "referrer": "https://staging.drivy.com/",
    "user_agent": "Mozilla/5.0 (Macintosh; Intel Mac OS X 10_1",
    "remote_encrypted_ip": "2a7c3399d5db02c55996365b99fd38d47a",
    "browser_width": "1440",
    "browser_height": "710",
    "browser_language": "en",
    "browser_timezone_offset": "+01:00",
    "browser_persistent_storage": "true",
    "anonymous_id": "59a718e6-6dab-44b7-98d5-27148cc43439",
    "user_id": "1252",
    "tracking_source": "frontend/beacon"
  }
}
```

Payload example



Client side Infrastructure

Captur - Data collection



Captur - Debug

Tracking debugger

Events

Query

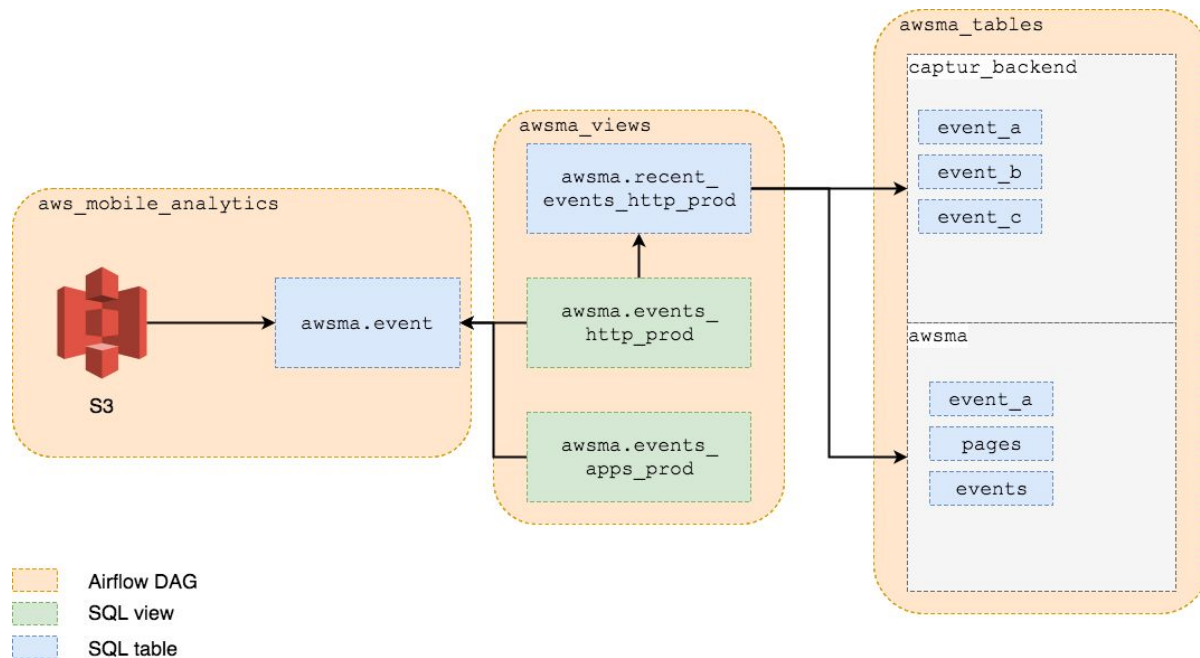
user_id

anonymous_id

[Clear screen](#)[Disable live polling](#)

Timestamp	Adapter	Operation	Payload
06/11/2018 17:21:57 CET	awsma	event	<pre>{ "type": "event", "attributes": { "rental_id": "4425082", "rental_state": "started", "user_role": "owner", "name": "rental_show_display_phone_number", "source": "mobile_webview", "anonymous_id": "unknown", "user_id": "2185628", "tracking_source": "backend" } }</pre>
06/11/2018 17:21:57 CET	awsma	page	<pre>{ "type": "page", "attributes": { "name": "homepage", "title": "Location voiture - Moin cher plus proche plus pratique - Drivy", "url": "https://www.drivy.com/", "path": "/", "user_agent": "Mozilla/5.0 (Macintosh; Intel Mac OS X 10_12_6) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/70.0.3538.77 Safari/537.36", "remote_ip": "62.23.191.0", "remote_encrypted_ip": "2a7c3399d5db02c55996365b99fd38647a07d083982bf8ecf22e108b2856ee05", "browser_width": "1440", "browser_height": "745", "browser_language": "en-G8", "browser_timezone_offset": "+01:00", "browser_persistent_storage": "true", "anonymous_id": "226ddffa-3741-4519-9eaa-44571e00cb2e", "user_id": "1839997", "navigation_type": "navigate", "tracking_source": "frontend/beacon" } }</pre>
06/11/2018 17:21:56 CET	awsma	event	<pre>{ "type": "event", "attributes": { "name": "unknownLate_address"</pre>

Captur: ETL side



Usage example: Tableau

Web Analytics

Ex: Booking Form Funnel

Usage example: Fraud Prevention



Drivy Rangers

Rules engine that can detect various types of fraud in minutes

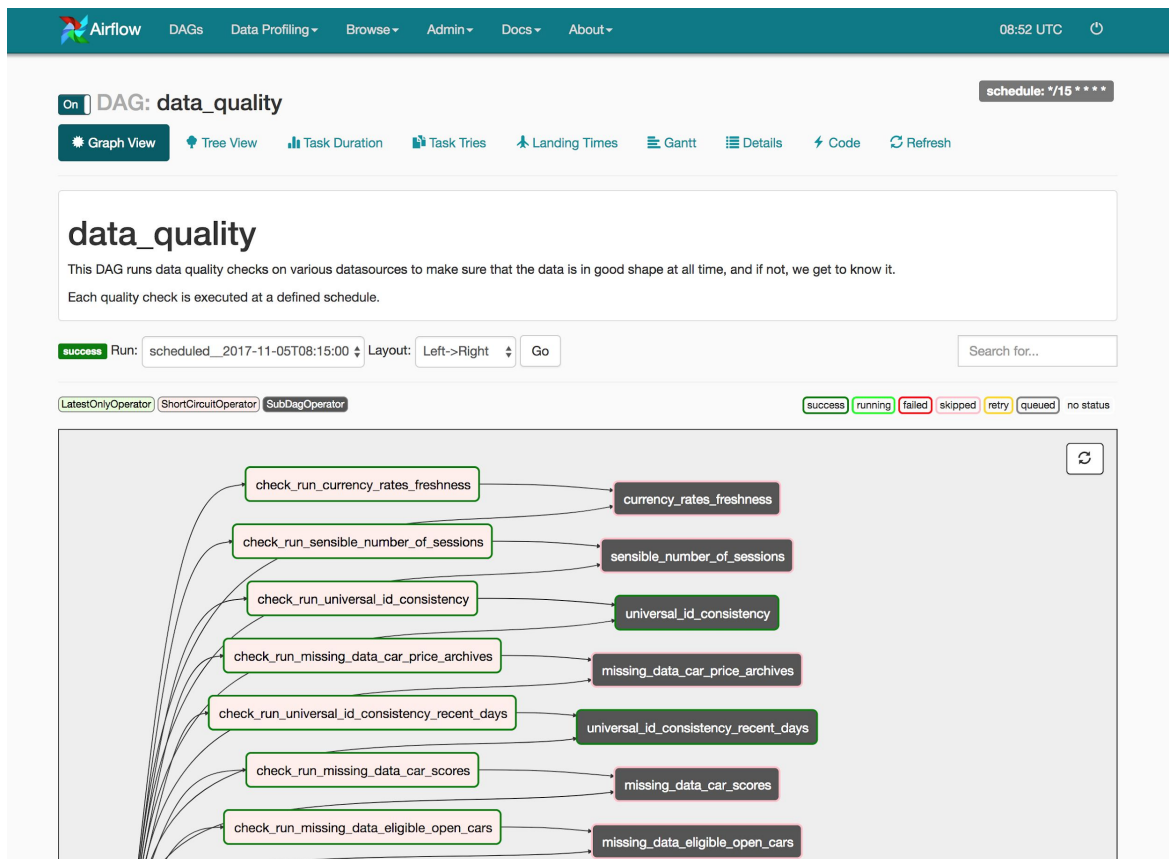
- >50 alerts in place (theft, insurance fraud, bad usage, phishing)
- >100 todos / day
- >40 users blacklisted / day
- Alerts added instantaneously in prod
- Escalation process per alerts (slack, pager duty...)
- Operated by 3 analysts 24/7

Data Quality Checkers


“Data quality checks are performed at a specified interval on one or multiple datasets that are coming from various datasources, using predicates we define. Checks have a tolerance and trigger alerts on alert destinations with an alert level defined by the severity of the found errors.”

Data Quality Checkers

- **ShortCircuitOperator**: determine if the checker should be executed or not.
- **SubDAGOperator**: actually runs the checks.



Data Quality Checker example

 Airflow

DAGs

Data Profiling ▾


Browse ▾

Admin ▾

Docs ▾

About ▾


08:53 UTC





SUBDAG: data_quality.sensible_number_of_sessions


schedule: */15 ****


[← Back to data_quality](#)


 **Graph View**


 Tree View


 Task Duration


 Task Tries

 Landing Times

 Gantt

 Details

 Code

 Refresh

Run: backfill_2017-11-04T15:45:00 ▾

Layout: Left->Right ▾

Go

Search for...

BashOperator

DummyOperator

PythonOperator

success

running


failed

skipped

retry


queued

no status



```
graph LR; A[extract_sensible_number_of_sessions_main] --> C[run_checker_sensible_number_of_sessions]; B[extract_sensible_number_of_sessions_other] --> C; C --> D[clean_tmp_sensible_number_of_sessions]
```

Questions?



Thank you!

